AWS Fundamentals course

2017, November, 14 1:13 PM

Application Load balancer:
- Static or Elastic IPs

Auto Scaling
1) Launch Config
2) Auto Scaling Group
    a. Contains multiple EC2 instances
        i. If one fails replace it
    b. Monitored using Amazon Cloud Watch

3) Optional
    a. Scheduled Action
    b. Scaling Policy

Cloud watch alarm examples
1) CPU > 80% for 5 periods of 1 min
    a. Add 100% (Scaling policies)
2) CPU < 40% for 5 periods of 1 minute
    a. Remove 50%

To allow applications to spin up new machines elastically
**Don't store state locally
**Don't rely on fixed IP addresses
**Don't use sticky sessions

Security:

Shared Responsibility Model
Resource Based Permissions
AWS CloudTrail
AWS Key Management Service
- Does not store encrypted data key
    - Needs to be handles by application

AWS Certificate Manager

***** DO NOT EMBED KEYS INTO CODE*****
- Instead utilize roles which are attached to EC2 instances and behind the scenes the AWS Token Service will request a Token and authenticate based on what we are trying to do
- Roles atuomatically last for 15 minutes and will auto rene
- Role is applied to EC2 instances when it is created and any application running on that EC2 instance will have access to the role