

Kafka Bootcamp

2017, December, 14 9:04 AM

Message Keys are what determines order

- Not required, producer can provide a NULL key which round robin the messages to partitions

Gwen Shapira

- ETL Lens
 - o Kafka is an integration platform not a point to point process
- Scalability
- Publisher/subscriber architecture, storage, processing
- Future?
 - o Additional languages (mainly java now)
 - o More security (PII etc)
 - o Cloud
- How do you know how many topics/partitions
 - o Think of topics as tables in a database
 - More flexible
 - Can change schema etc...
 - o # of partitions is chosen on scale
- Data should be stored in a raw format, kafka has an immutable datastore
 - o History has to be repeatable
 - o Audit purposes
- Data has to be encrypted by consumer

Chris Matta

- Why Streaming?
 - o All data is even streaming and has been for some time
 - o Databases are materializations of events, Change Data Capture
 - o Capture allows these CDC logs to be centralized
- Apache Kafka
 - o Built to handle data flowing into Hadoop
 - Needed a system to act as a shock absorber for the volume of data into HDFS
 - "Immutable append only distributed commit log"
 - Distributed by design
 - ◆ Replication
 - ◆ Fault Tolerance
 - ◆ Partitioning
 - ◆ Elastic Scaling
 - Cannot edit messages in queue and cannot delete an individual message (except for compact queue??)
 - o Runs on commodity linux host
 - o Strict ordering and guaranteed write
 - Persistence
 - Consumers job to acknowledge it read the message then the partition deletes it
 - o When you write data into Kafka you write it into memory of the broker
 - Data must be serialized as bytes so you need to know the data when you deserialize it
 - Kafka says this topic offset points to this memory offset in unix
 - ◆ Java writes to unix memory log
 - ◆ Linux page cache is very important
 - ◇ Linux page cache is held in memory for a certain amount of time then flushed to disk
 - ◇ When data is read it will need to be promoted to cache as well
 - ▶ Can cause instability when consumers are reading slowly
 - o While you are writing data consumers can be reading from any point in the log
 - o If you provide no key then round robin is used to store values to topics

- If you do provide a key then the key is hashed to determine which partition it is written too
- When you write to Kafka you are writing to the leader then it is mirrored to the other boxes
- In the event of a failure a new leader is chosen from the topics that are in sync
- Kafka will "self heal" and will try to move leaders so that one broker is not a leader for more than one topic
- Order is guaranteed per partition but NOT along partitions
- Processing of data and writing/reading of data is decoupled
- Bootstrap URL is used to introduce the produce to each of the brokers provided in the bootstrap
- "Topics are metadata that wrap partitions"
- Want to plan the # of partitions that you can scale into so that you can horizontally scale partitions
 - If you have 5 servers and 4 partitions it is hard to setup a 5th partitions
 - But if you have 8 partitions and 4 servers it is relatively easy to scale up brokers
- Producers can set a config to when the brokers send a notification that messages are ready
 - Ack 0
 - No acknowledgement
 - Fastest way to write to kafka
 - Ack 1
 - Acknowledge that the broker wrote the message to itself
 - Ack all (-1)
 - Min.insync.replica=2
- Every 3 minutes consumer will send a message to consumer offset topic saying where it is at in the queue
 - This topic is typically replicated across the cluster
- Consumers
 - Consumers give themselves a name (consumer groups)
 - Multiple consumers can be in the consumer group, and can be given their own partition to read messages
 - Heartbeat for health of each consumer
 - This also gives us the ability to rebalance if a consumer dies
 - If you have more consumers than partitions then a consumer will sit idle
 - Can be written to consume from 1 partition
- Kafka Connect
 - Piece of the Kafka toolbox
 - Noticed lots of boiler print kafka producers and consumers
 - Instead you can download a Kafka connectors and configure it
 - Classified as sources or Sinks
- Kafka Streams
 - 3 paradigms of programming
 - Request processing
 - Batch processing
 - Stream processing
 - Event driven micro service
 - Simple java library
 - Look into creating a simple kafka stream app
 - ** No micro batch, true event processing **
 - Null keys are not supported in Kafka streams
 - Compacted topic will retire records where we have a new record for that key
 - Can read and write from topics
 - Ktable is a compacted topic in Kafka

***** # of partitions is the amount of parallelism*****

"One is none and 2 is one"

***** Keys will Determine Partition *****

***** avro is a serialization schema *****

Things to go over:

- Load balancing
- Change Data capture

- Z os linux OS
- Throughput of data
- Serializing data
- Linux page cache and linux heap and cache miss/hit
- Zookeeper
- <https://www.confluent.io/resources>
- Eclipse neon
- <https://github.com/cjmatta?tab=repositories>
- STAR schema